

```
1 %pip install numpy pandas river
```

```
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Collecting river
  Downloading river-0.23.0-cp312-cp312-manylinux_2_28_x86_64.whl.metadata (9.2 kB)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Collecting pandas
  Downloading pandas-2.3.3-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl.metadata (91 kB)
    91.2/91.2 kB 3.5 MB/s eta 0:00:00
Requirement already satisfied: scipy<2.0.0,>=1.14.1 in /usr/local/lib/python3.12/dist-packages (from river) (1.16.3)
Requirement already satisfied: six>1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
  Downloading river-0.23.0-cp312-cp312-manylinux_2_28_x86_64.whl (3.2 MB)
    3.2/3.2 MB 48.1 MB/s eta 0:00:00
  Downloading pandas-2.3.3-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl (12.4 MB)
    12.4/12.4 MB 116.6 MB/s eta 0:00:00
Installing collected packages: pandas, river
  Attempting uninstall: pandas
    Found existing installation: pandas 2.2.2
    Uninstalling pandas-2.2.2:
      Successfully uninstalled pandas-2.2.2
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the
google-colab 1.0.0 requires pandas==2.2.2, but you have pandas 2.3.3 which is incompatible.
Successfully installed pandas-2.3.3 river-0.23.0
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import kagglehub
6 import time
```

```
1 path = kagglehub.dataset_download("mlg-ulb/creditcardfraud")
2
3 print("Path to dataset files:", path)
```

Using Colab cache for faster access to the 'creditcardfraud' dataset.  
Path to dataset files: /kaggle/input/creditcardfraud

```
1 df = pd.read_csv(path+'/creditcard.csv')
2 df.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.11
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.10
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.90
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.19
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.13

5 rows × 31 columns

```
1 df.describe()
```

	Time	V1	V2	V3	V4	V5	V6	V7
count	284807.000000	2.848070e+05						
std	47.485145055	9.58896e-001	1.851509e-001	1.516255e-001	1.415809e-001	1.380274e-001	1.352271e-001	1.237094e-001
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01

This notebook evaluates various online machine learning models from the River library for credit card fraud detection. We'll compare their performance using F1 score and MCC, considering different scaling and sampling techniques, and visualize their historical performance over time.

```

1 import river
2 from river.linear_model import LogisticRegression
3 from river.forest import ARFClassifier
4 from river.linear_model import PAClassifier
5 from river.ensemble import AdaBoostClassifier, ADWINBoostingClassifier
6 from river.tree import HoeffdingTreeClassifier
7 from river.imblearn import HardSamplingClassifier, RandomSampler
8 from river.metrics import F1, MCC
9 from river.preprocessing import StandardScaler
10 from river.stream import iter_pandas

```

1 Start coding or generate with AI.

## Modulation

Creating a Python function `evaluate_model_pipeline` that takes a River model, its name, and the feature/target data ( $X, y$ ) as input. This function will encapsulate the online learning loop, calculate F1 and MCC scores, print progress, and return the final F1 and MCC scores along with runtime.

Then, preparing a list of the following River models (including their corresponding scalers and sampling techniques, if applicable) that have been used previously in the notebook:

1. Logistic Regression (No Scaling)
2. Scaled Logistic Regression
3. Scaled Adaptive Random Forest (`ARFClassifier`)
4. Scaled Passive Aggressive Classifier (`PAClassifier`)
5. Scaled AdaBoost (with `HoeffdingTreeClassifier`)
6. `ARFClassifier` with HardSampling
7. AdaBoost with HardSampling (with Scaled `HoeffdingTreeClassifier`)
8. `ADWINBoostingClassifier` (with Scaled `LogisticRegression`)

Iterating through this list, call the `evaluate_model_pipeline` function for each model, and store the returned F1, MCC, and ROC AUC scores along with the model name. Finally, display all the collected results in a Pandas DataFrame and generate bar plots for F1 Score and MCC Score to visually compare the models.

## Evaluate Model Function

```

1 # @title Evaluate Model Function
2
3 X_data = df.drop('Class', axis=1)
4 y_data = df['Class']
5
6 def evaluate_model_pipeline(model, model_name, X, y):
7     """
8         Evaluates a River model pipeline using online learning, calculating F1 and MCC scores.
9         Stores historical scores at 10,000-row intervals.
10
11     Args:
12         model: The River model pipeline (e.g., StandardScaler() | LogisticRegression()).
13         model_name (str): The name of the model for printing progress.
14         X (pd.DataFrame): The feature DataFrame.
15         y (pd.Series): The target Series.
16
17     Returns:
18         tuple: A tuple containing the final F1 score, final MCC score, lists of historical F1 scores, historical MCC scores.
19     """
20
21     # Instantiate the metrics
22     f1_metric = F1()

```

```

23     mcc_metric = MCC()
24
25     # Initialize empty lists to store historical scores and row numbers
26     history_f1 = []
27     history_mcc = []
28     row_numbers = []
29
30
31     # Start timing
32     start_time = time.time()
33
34     # Iterate through the DataFrame instance by instance
35     for i, (x, y_true) in enumerate(iter_pandas(X, y)):
36
37         # Predict and update metrics
38         y_pred = model.predict_one(x)
39
40         f1_metric.update(y_true, y_pred)
41         mcc_metric.update(y_true, y_pred)
42
43         # Learn from the true label
44         model.learn_one(x, y_true)
45
46         # Store scores every 10000 rows
47         if (i + 1) % 10000 == 0:
48             history_f1.append(f1_metric.get())
49             history_mcc.append(mcc_metric.get())
50             row_numbers.append(i + 1)
51
52
53     # End timing
54     end_time = time.time()
55     total_time = end_time - start_time
56
57     # Print final scores
58     final_f1 = f1_metric.get()
59     final_mcc = mcc_metric.get()
60     print(f"\nFinal Scores with {model_name}:")
61     print(f"Final F1 Score = {final_f1:.4f}")
62     print(f"Final MCC Score = {final_mcc:.4f}")
63     print(f"Total time taken: {total_time:.2f} seconds")
64

```

## ▼ Logistic Regression Models

Logistic Regression is a linear model for binary classification. We evaluate it with and without feature scaling, and also with HardSampling to address class imbalance.

## ▼ Logistic Regression Models

```

1 # @title Logistic Regression Models
2
3 # Evaluate Logistic Regression (No Scaling)
4 model_ns_logreg = LogisticRegression()
5 final_f1_ns, final_mcc_ns, hist_f1_ns, hist_mcc_ns, rows_ns, time_ns = evaluate_model_pipeline(model_ns_logreg, 'Logistic
6
7
8 # Evaluate Scaled Logistic Regression
9 model_sc_logreg = StandardScaler() | LogisticRegression()
10 final_f1_sc, final_mcc_sc, hist_f1_sc, hist_mcc_sc, rows_sc, time_sc = evaluate_model_pipeline(model_sc_logreg, 'Scaled L
11
12
13
14 # Evaluate HardSampling Logistic Regression
15 model_hs_logreg = HardSamplingClassifier(classifier=model_sc_logreg, size=100, p=0.2, seed=42)
16 final_f1_hs_logreg, final_mcc_hs_logreg, hist_f1_hs_logreg, hist_mcc_hs_logreg, rows_hs_logreg, time_hs_logreg = evaluate_
17
18
19

```

```

20 # Store results for final DataFrame
21 results_summary = [
22     {'Model': 'Logistic Regression (No Scaling)', 'Final F1 Score': final_f1_ns, 'Final MCC Score': final_mcc_ns, 'Total Runtime': final_time_ns, 'Rows': final_rows_ns, 'Time per Row': final_time_per_row_ns},
23     {'Model': 'Scaled Logistic Regression', 'Final F1 Score': final_f1_sc, 'Final MCC Score': final_mcc_sc, 'Total Runtime': final_time_sc, 'Rows': final_rows_sc, 'Time per Row': final_time_per_row_sc},
24     {'Model': 'HardSampling Logistic Regression', 'Final F1 Score': final_f1_hs_logreg, 'Final MCC Score': final_mcc_hs_logreg, 'Total Runtime': final_time_hs_logreg, 'Rows': final_rows_hs_logreg, 'Time per Row': final_time_per_row_hs_logreg}
25 ]
26
27 df_summary = pd.DataFrame(results_summary)
28 print("\n--- Final Scores Summary ---")
29

```

```

Final Scores with Logistic Regression (No Scaling):
Final F1 Score = 0.0588
Final MCC Score = 0.0572
Total time taken: 14.95 seconds

Final Scores with Scaled Logistic Regression:
Final F1 Score = 0.7933
Final MCC Score = 0.7969
Total time taken: 34.13 seconds

Final Scores with HardSampling Logistic Regression:
Final F1 Score = 0.4414
Final MCC Score = 0.4525
Total time taken: 40.93 seconds

--- Final Scores Summary ---
      Model  Final F1 Score  Final MCC Score \
0  Logistic Regression (No Scaling)      0.058824      0.057192
1    Scaled Logistic Regression      0.793296      0.796936
2  HardSampling Logistic Regression      0.441397      0.452487

   Total Runtime
0        14.948919
1       34.126688
2      40.930604

```

## Adaptive Random Forest Models

Adaptive Random Forest (ARF) is an ensemble method for online learning. It handles concept drift and works well with data streams. We evaluate it with and without scaling, and with HardSampling.

## Adaptive Random Forest Models

```

1 # @title Adaptive Random Forest Models
2
3 # Evaluate Adaptive Random Forest (ARF) (no Scaling)
4 model_ARF = ARFClassifier(n_models=10, seed=42)
5 final_f1_ARF, final_mcc_ARF, hist_f1_ARF, hist_mcc_ARF, rows_ARF, time_ARF = evaluate_model_pipeline(model_ARF, 'Adaptive Random Forest (No Scaling)')
6
7
8 # Evaluate Scaled Adaptive Random Forest (ARF)
9 model_sc_ARF = StandardScaler() | ARFClassifier(n_models=10, seed=42)
10 final_f1_sc_ARF, final_mcc_sc_ARF, hist_f1_sc_ARF, hist_mcc_sc_ARF, rows_sc_ARF, time_sc_ARF = evaluate_model_pipeline(model_sc_ARF, 'Adaptive Random Forest (Scaled)')
11
12
13 # Evaluate Hardsampling with Adaptive Random Forest
14 base_arf = ARFClassifier(n_models=10, seed=42)
15 model_hs_arf = HardSamplingClassifier(classifier=base_arf, size=100, p=0.2, seed=42)
16
17 final_f1_hs_ARF, final_mcc_hs_ARF, hist_f1_hs_ARF, hist_mcc_hs_ARF, rows_hs_ARF, time_hs_ARF = evaluate_model_pipeline(model_hs_arf, 'Adaptive Random Forest (Hardsampling)')
18
19
20

```

```

21 # Store results for the final dataframe
22 results_summary.extend([
23     {'Model': 'Adaptive Random Forest (no Scaling)', 'Final F1 Score': final_f1_ARF, 'Final MCC Score': final_mcc_ARF, 'Total time taken': total_time_ARF},
24     {'Model': 'Scaled Adaptive Random Forest', 'Final F1 Score': final_f1_sc_ARF, 'Final MCC Score': final_mcc_sc_ARF, 'Total time taken': total_time_sc_ARF},
25     {'Model': 'HardSampling Adaptive Random Forest', 'Final F1 Score': final_f1_hs_ARF, 'Final MCC Score': final_mcc_hs_ARF, 'Total time taken': total_time_hs_ARF}
26 ])
27
28 df_summary = pd.DataFrame(results_summary)
29 print("\n--- Final Scores Summary ---")

```

Final Scores with Adaptive Random Forest (no Scaling):  
 Final F1 Score = 0.5008  
 Final MCC Score = 0.5720  
 Total time taken: 512.50 seconds

Final Scores with Scaled Adaptive Random Forest:  
 Final F1 Score = 0.4102  
 Final MCC Score = 0.4965  
 Total time taken: 460.02 seconds

Final Scores with HardSampling Adaptive Random Forest:  
 Final F1 Score = 0.7371  
 Final MCC Score = 0.7534  
 Total time taken: 565.80 seconds

--- Final Scores Summary ---

	Model	Final F1 Score	Final MCC Score	\
3	Adaptive Random Forest (no Scaling)	0.500754	0.571952	
4	Scaled Adaptive Random Forest	0.410175	0.496511	
5	HardSampling Adaptive Random Forest	0.737101	0.753400	
Total Runtime				
3	512.502507			
4	460.019537			
5	565.795731			

## ▼ Passive Aggressive Classifier Models

Passive Aggressive Classifier (PAClassifier) is a simple yet effective online learning algorithm that updates the model only when a misclassification occurs. We evaluate it with and without scaling, and with HardSampling.

## ▼ PassiveAgressive Classifier Models

```

1 # @title PassiveAgressive Classifier Models
2
3 #Evaluate PassiveAggressive Classifier (no Scaling)
4 model_pa = PAClassifier(C=0.01, mode=1)
5 final_f1_pa, final_mcc_pa, hist_f1_pa, hist_mcc_pa, rows_pa, time_pa = evaluate_model_pipeline(model_pa, 'Passive Aggressive Classifier (no Scaling)')
6
7
8
9 # Evaluate Scaled PassiveAgressive Classifier
10 model_sc_pa = StandardScaler() | PAClassifier(C=0.01, mode=1)
11 final_f1_sc_pa, final_mcc_sc_pa, hist_f1_sc_pa, hist_mcc_sc_pa, rows_sc_pa, time_sc_pa = evaluate_model_pipeline(model_sc_pa, 'Passive Aggressive Classifier (Scaled)')
12
13
14
15 # Evaluate HardSampling with PassiveAgressive Classifiers
16 base_pa = StandardScaler() | PAClassifier(C=0.01, mode=1)
17 model_hs_pa = HardSamplingClassifier(classifier=base_pa, size=100, p=0.2, seed=42)
18
19 final_f1_hs_pa, final_mcc_hs_pa, hist_f1_hs_pa, hist_mcc_hs_pa, rows_hs_pa, time_hs_pa = evaluate_model_pipeline(model_hs_pa, 'HardSampling Adaptive Random Forest')
20

```

```

21
22
23 # Store results for the final dataframe
24 results_summary.extend([
25     {'Model': 'Passive Aggressive Classifier (no Scaling)', 'Final F1 Score': final_f1_pa, 'Final MCC Score': final_mcc_pa,
26     {'Model': 'Scaled Passive Aggressive Classifier', 'Final F1 Score': final_f1_sc_pa, 'Final MCC Score': final_mcc_sc_pa,
27     {'Model': 'HardSampling Passive Aggressive Classifier', 'Final F1 Score': final_f1_hs_pa, 'Final MCC Score': final_mcc_hs_pa,
28 }]
29
30 df_summary = pd.DataFrame(results_summary)
31 print("\n--- Final Scores Summary ---")
32 print(df_summary.iloc[6:9])

```

Final Scores with Passive Aggressive Classifier (no Scaling):  
Final F1 Score = 0.0528  
Final MCC Score = 0.0512  
Total time taken: 11.48 seconds

Final Scores with Scaled Passive Aggressive Classifier:  
Final F1 Score = 0.8000  
Final MCC Score = 0.8020  
Total time taken: 28.52 seconds

Final Scores with HardSampling Passive Aggressive Classifier:  
Final F1 Score = 0.5483  
Final MCC Score = 0.5972  
Total time taken: 39.36 seconds

--- Final Scores Summary ---

	Model	Final F1 Score	Final MCC Score
6	Passive Aggressive Classifier (no Scaling)	0.052830	0.051206
7	Scaled Passive Aggressive Classifier	0.800000	0.801971
8	HardSampling Passive Aggressive Classifier	0.548295	0.597199

Total Runtime

6	11.480768
7	28.523685
8	39.362953

## ▼ AdaBoost Classifier Models

AdaBoost (Adaptive Boosting) is an ensemble meta-algorithm that can be used with a variety of other learning algorithms to improve performance. We use it with Logistic Regression, Hoeffding Trees, and Passive Aggressive Classifiers as base estimators.

## ▼ AdaBoost Classifier Models

```

1 # @title AdaBoost Classifier Models
2
3 # Evaluate AdaBoost Logistic Regression
4 model_ada = AdaBoostClassifier(model=StandardScaler() | LogisticRegression(), n_models = 10, seed=42)
5 final_f1_ada, final_mcc_ada, hist_f1_ada, hist_mcc_ada, rows_ada, time_ada = evaluate_model_pipeline(model_ada, 'AdaBoost')
6
7
8 # Evaluate AdaBoost PassiveAggressive Classifier
9 model_ada_pa = AdaBoostClassifier(model=StandardScaler() | PAClassifier(C=0.1, mode=1), n_models = 10, seed=42)
10 final_f1_ada_pa, final_mcc_ada_pa, hist_f1_ada_pa, hist_mcc_ada_pa, rows_ada_pa, time_ada_pa = evaluate_model_pipeline(model_ada_pa)
11
12
13 # Store scores in the result summary
14 results_summary.extend([
15     {'Model': 'AdaBoost Logistic Regression', 'Final F1 Score': final_f1_ada, 'Final MCC Score': final_mcc_ada, 'Total Run time': time_ada},
16     {'Model': 'AdaBoost PassiveAggressive Classifier', 'Final F1 Score': final_f1_ada_pa, 'Final MCC Score': final_mcc_ada_pa, 'Total Run time': time_ada_pa}
17 ])

```

```
18
19 df_summary = pd.DataFrame(results_summary)
20 print("\n--- Final Scores Summary ---")
21
```

```
Final Scores with AdaBoost Logistic Regression:
Final F1 Score = 0.7404
Final MCC Score = 0.7400
Total time taken: 246.54 seconds
```

```
Final Scores with AdaBoost PassiveAgressive Classifier:
Final F1 Score = 0.7518
Final MCC Score = 0.7514
Total time taken: 216.25 seconds
```

```
--- Final Scores Summary ---
Model   Final F1 Score  Final MCC Score \
9       AdaBoost Logistic Regression      0.740443    0.740027
10     AdaBoost PassiveAgressive Classifier      0.751788    0.751370

Total Runtime
9        246.536836
10      216.249747
```

```
1 # Evaluate AdaBoost Hoeffding Tree
2 model_ada_ht = AdaBoostClassifier(model=StandardScaler() | HoeffdingTreeClassifier(), n_models = 10, seed=42)
3 final_f1_ada_ht, final_mcc_ada_ht, hist_f1_ada_ht, hist_mcc_ada_ht, rows_ada_ht, time_ada_ht = evaluate_model_pipeline(mode
4
5 results_summary.extend([
6     {'Model': 'AdaBoost Hoeffding Tree', 'Final F1 Score': final_f1_ada_ht, 'Final MCC Score': final_mcc_ada_ht, 'Total Rur
7 ])
8
9 df_summary = pd.DataFrame(results_summary)
10 print("\n--- Final Scores Summary ---")
11 print(df_summary.iloc[11:12])
```

```
Final Scores with AdaBoost Hoeffding Tree:
Final F1 Score = 0.7903
Final MCC Score = 0.7929
Total time taken: 972.22 seconds
```

```
--- Final Scores Summary ---
Model   Final F1 Score  Final MCC Score  Total Runtime
11 AdaBoost Hoeffding Tree      0.790287      0.792909    972.216348
```

1 Start coding or generate with AI.

## ✓ ADWINBoosting Classifier Models

ADWINBoostingClassifier is an ensemble method that combines AdaBoost with the Adaptive Windowing (ADWIN) algorithm, allowing it to adapt to concept drift. We evaluate it with Logistic Regression and Passive Aggressive Classifiers as base estimators.

```
1 # Evaluate ADWINBoosting with Logistic Regression
2 model_adwin_lr = ADWINBoostingClassifier(model=StandardScaler() | LogisticRegression(), n_models=10,seed=42)
3 final_f1_adwin_lr, final_mcc_adwin_lr, hist_f1_adwin_lr, hist_mcc_adwin_lr, rows_adwin_lr, time_adwin_lr = evaluate_model_
4
5 # Evaluate ADWINBoosting with PassiveAgressive Classifier
6 model_adwin_pa = ADWINBoostingClassifier(model=StandardScaler() | PAClassifier(C=0.1, mode=1), n_models=10,seed=42)
7 final_f1_adwin_pa, final_mcc_adwin_pa, hist_f1_adwin_pa, hist_mcc_adwin_pa, rows_adwin_pa, time_adwin_pa = evaluate_model_
8
9
```

```

10 # Store the results
11 results_summary.extend([
12     {'Model': 'ADWINBoosting Logistic Regression', 'Final F1 Score': final_f1_adwin_lr, 'Final MCC Score': final_mcc_adwin_lr},
13     {'Model': 'ADWINBoosting PassiveAggressive Classifier', 'Final F1 Score': final_f1_adwin_pa, 'Final MCC Score': final_mcc_adwin_pa},
14 ])
15
16 df_summary = pd.DataFrame(results_summary)
17 print("\n--- Final Scores Summary ---")
18 print(df_summary)

```

Final Scores with ADWINBoosting Logistic Regression:

Final F1 Score = 0.7955

Final MCC Score = 0.7997

Total time taken: 427.97 seconds

Final Scores with ADWINBoosting PassiveAggressive Classifier:

Final F1 Score = 0.7893

Final MCC Score = 0.7900

Total time taken: 430.26 seconds

--- Final Scores Summary ---

	Model	Final F1 Score	\
12	ADWINBoosting Logistic Regression	0.795506	
13	ADWINBoosting PassiveAggressive Classifier	0.789305	
	Final MCC Score	Total Runtime	
12	0.799676	427.972467	
13	0.790049	430.263999	

1 Start coding or [generate](#) with AI.

1 Start coding or [generate](#) with AI.

## ▼ Results

1 df\_summary.round(4)

	Model	Final F1 Score	Final MCC Score	Total Runtime		
0	Logistic Regression (No Scaling)	0.0588	0.0572	14.9489		
1	Scaled Logistic Regression	0.7933	0.7969	34.1267		
2	HardSampling Logistic Regression	0.4414	0.4525	40.9306		
3	Adaptive Random Forest (no Scaling)	0.5008	0.5720	512.5025		
4	Scaled Adaptive Random Forest	0.4102	0.4965	460.0195		
5	HardSampling Adaptive Random Forest	0.7371	0.7534	565.7957		
6	Passive Aggressive Classifier (no Scaling)	0.0528	0.0512	11.4808		
7	Scaled Passive Aggressive Classifier	0.8000	0.8020	28.5237		
8	HardSampling Passive Aggressive Classifier	0.5483	0.5972	39.3630		
9	AdaBoost Logistic Regression	0.7404	0.7400	246.5368		
10	AdaBoost PassiveAggressive Classifier	0.7518	0.7514	216.2497		
11	AdaBoost Hoeffding Tree	0.7903	0.7929	972.2163		
12	ADWINBoosting Logistic Regression	0.7955	0.7997	427.9725		
13	ADWINBoosting PassiveAggressive Classifier	0.7893	0.7900	430.2640		

1 # Set the index to model names for easier plotting

```

2 df_results = df_summary.set_index('Model')
3
4 # Plotting the scores
5 df_f1_sorted = df_results.sort_values(by='Final F1 Score', ascending=False)
6
7 # F1 Score plot
8 plt.figure(figsize=(12, 10))
9 sns.barplot(x=df_f1_sorted.index, y='Final F1 Score', data=df_f1_sorted,
10 palette='magma')
11 plt.title('Comparison of Model Final F1 Scores (Sorted Descending)')
12 plt.xlabel('Model')
13 plt.ylabel('F1 Score')
14 plt.xticks(rotation=90, ha='right')
15 plt.tight_layout()
16 plt.show()
17

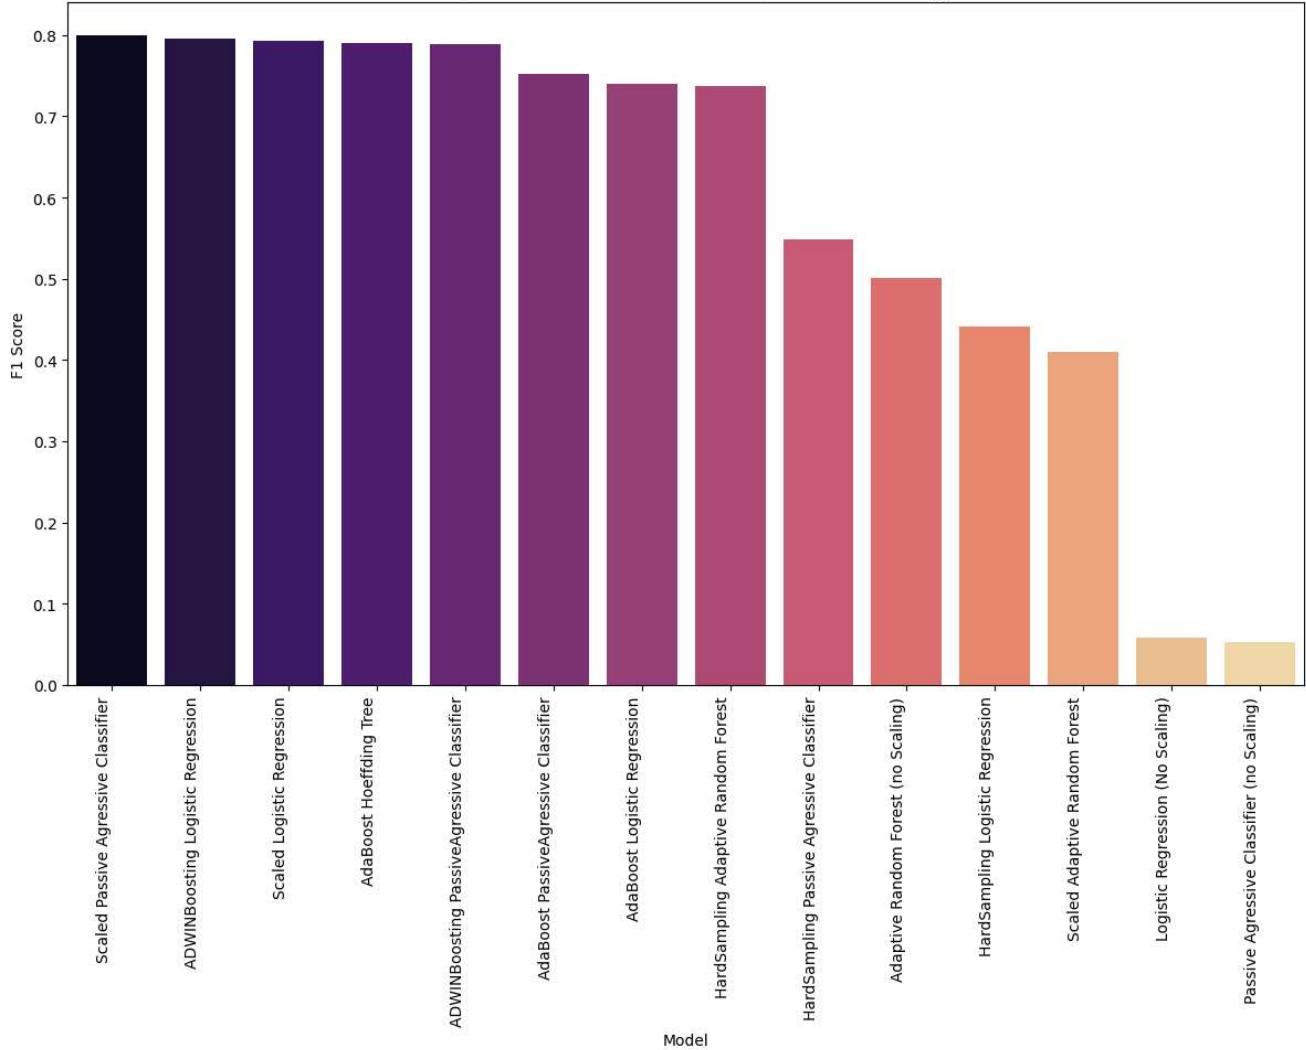
```

/tmp/ipython-input-3445809527.py:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set

```
sns.barplot(x=df_f1_sorted.index, y='Final F1 Score', data=df_f1_sorted, palette='magma')
```

Comparison of Model Final F1 Scores (Sorted Descending)



```

1 df_mcc_sorted = df_results.sort_values(by='Final MCC Score', ascending=False)
2
3 # MCC Score plot
4 plt.figure(figsize=(12, 10))
5 sns.barplot(x=df_mcc_sorted.index, y='Final MCC Score', data=df_mcc_sorted, palette='magma')
6 plt.title('Comparison of Model Final MCC Scores (Sorted Descending)')
7 plt.xlabel('Model')
8 plt.ylabel('MCC Score')
9 plt.xticks(rotation=90, ha='right')
10 plt.tight_layout()
11 plt.show()

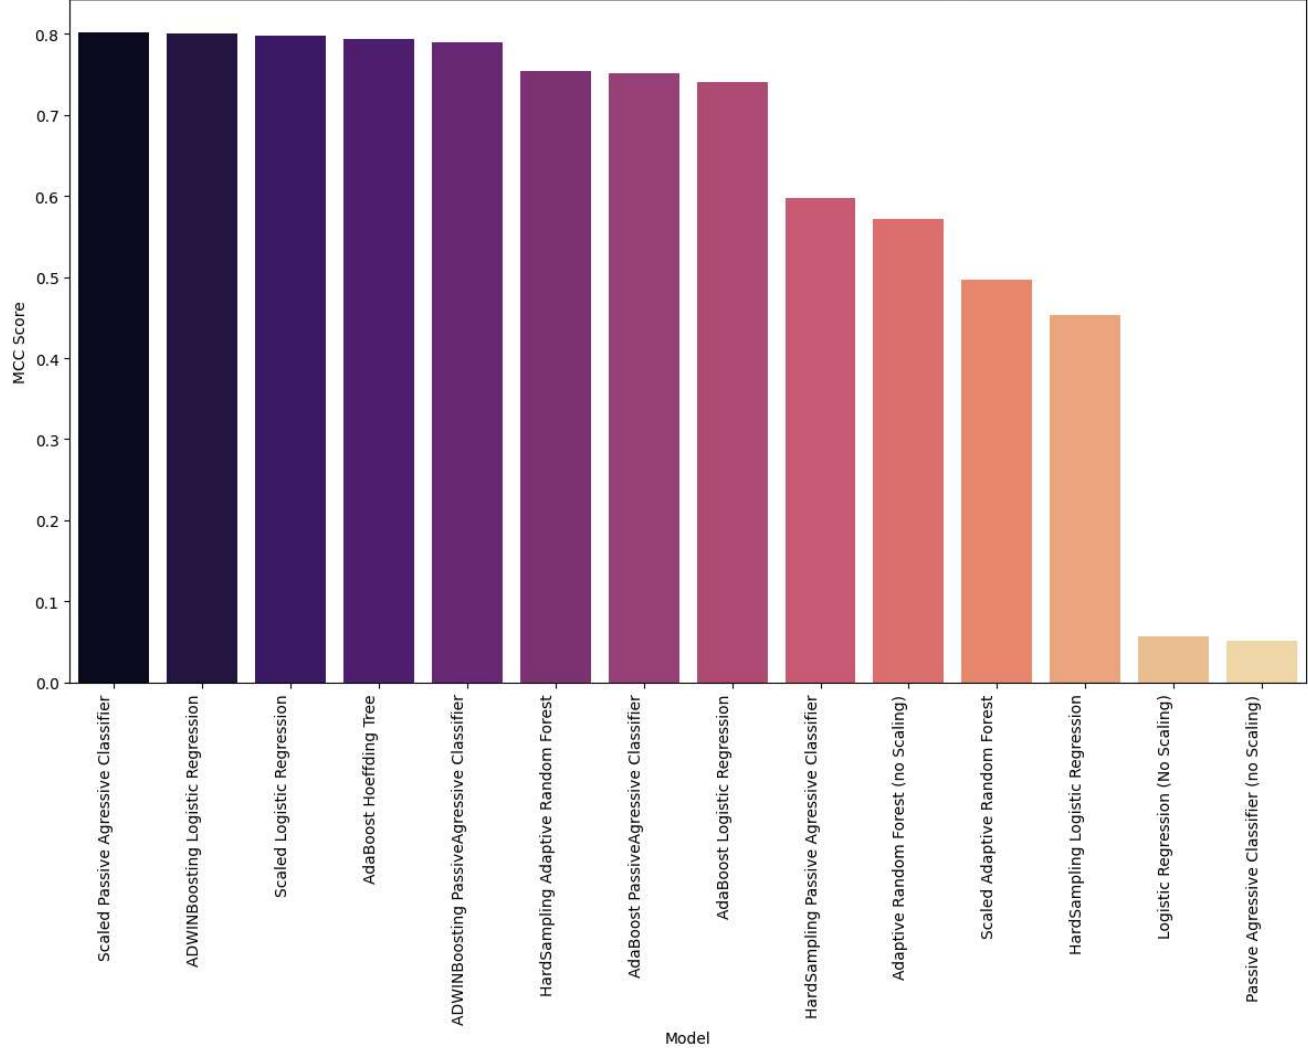
```

/tmp/ipython-input-3174712153.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set

```
sns.barplot(x=df_mcc_sorted.index, y='Final MCC Score', data=df_mcc_sorted, palette='magma')
```

Comparison of Model Final MCC Scores (Sorted Descending)



```

1 df_runtime_sorted = df_results.sort_values(by='Total Runtime', ascending=True)
2
3 # Total Runtime plot
4 plt.figure(figsize=(12, 10))
5 sns.barplot(x=df_runtime_sorted.index, y='Total Runtime', data=df_runtime_sorted, palette='magma')

```

```

6 plt.title('Comparison of Model Total Runtimes (Sorted Descending)')
7 plt.xlabel('Model')
8 plt.ylabel('Runtime (seconds)')
9 plt.xticks(rotation=90, ha='right')
10 plt.tight_layout()

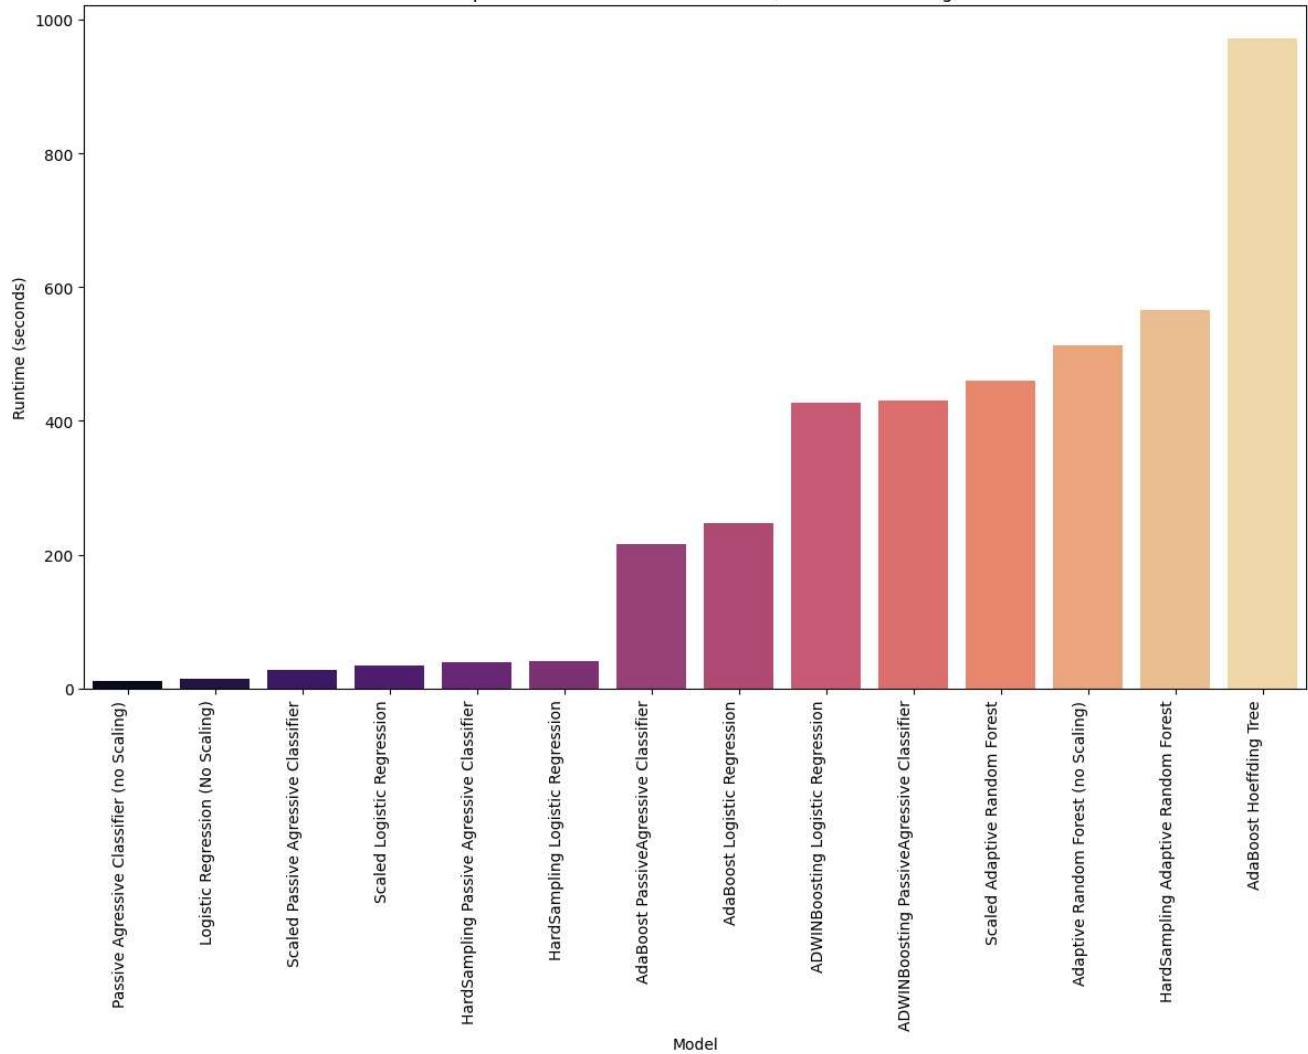
```

/tmp/ipython-input-3676103878.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set

```
sns.barplot(x=df_runtime_sorted.index, y='Total Runtime', data=df_runtime_sorted, palette='magma')
```

Comparison of Model Total Runtimes (Sorted Descending)



```

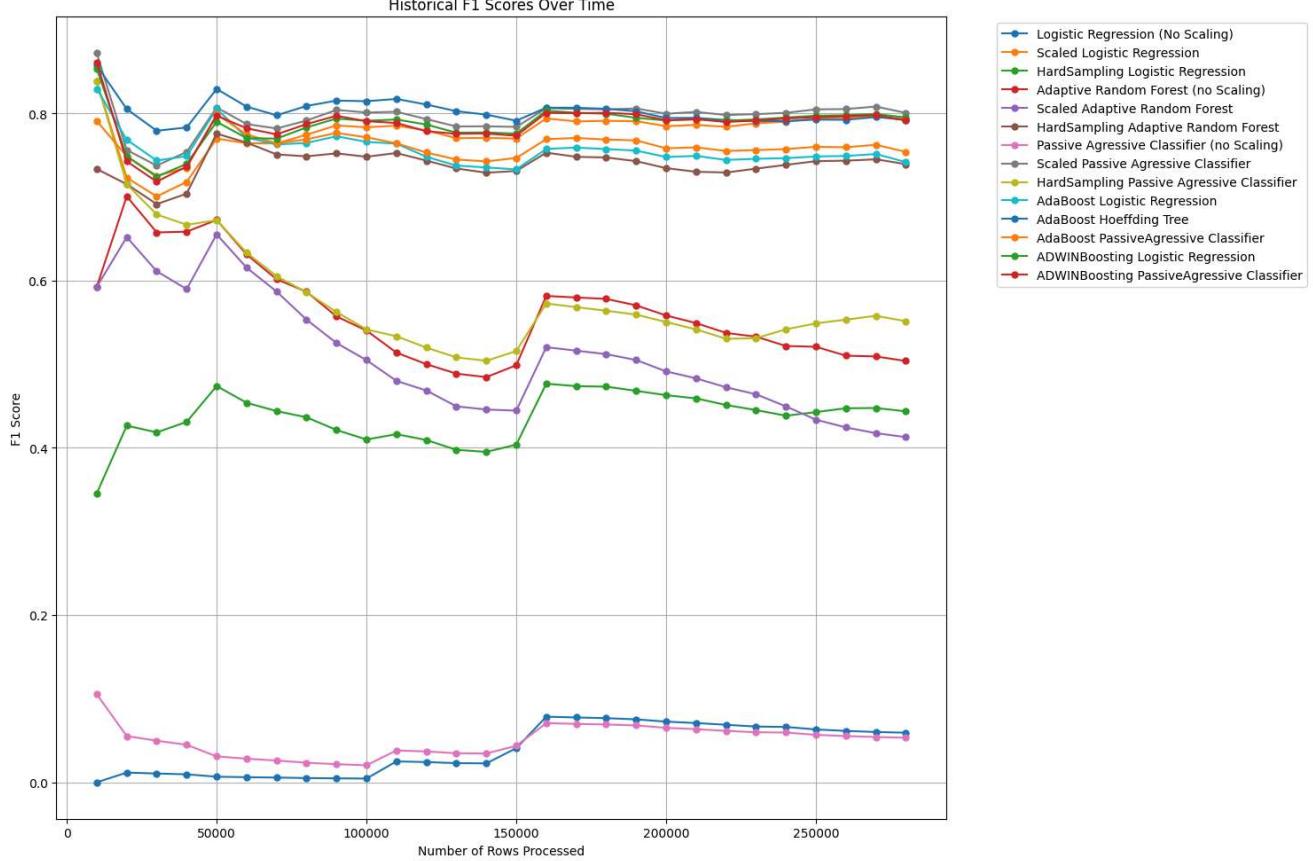
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 # Collect all historical data
5 historical_data = []
6
7 historical_data.extend([
8     {'Model': 'Logistic Regression (No Scaling)', 'f1_history': hist_f1_ns, 'mcc_history': hist_mcc_ns, 'rows': rows_ns},
9     {'Model': 'Scaled Logistic Regression', 'f1_history': hist_f1_sc, 'mcc_history': hist_mcc_sc, 'rows': rows_sc},
10    {'Model': 'HardSampling Logistic Regression', 'f1_history': hist_f1_hs_logreg, 'mcc_history': hist_mcc_hs_logreg, 'rows': rows_hs_logreg},
11    {'Model': 'Adaptive Random Forest (no Scaling)', 'f1_history': hist_f1_ARF, 'mcc_history': hist_mcc_ARF, 'rows': rows_ARF}
])

```

```

12     {'Model': 'Scaled Adaptive Random Forest', 'f1_history': hist_f1_sc_ARF, 'mcc_history': hist_mcc_sc_ARF, 'rows': rows_
13     {'Model': 'HardSampling Adaptive Random Forest', 'f1_history': hist_f1_hs_ARF, 'mcc_history': hist_mcc_hs_ARF, 'rows'_
14     {'Model': 'Passive Agressive Classifier (no Scaling)', 'f1_history': hist_f1_pa, 'mcc_history': hist_mcc_pa, 'rows': i_
15     {'Model': 'Scaled Passive Aggressive Classifier', 'f1_history': hist_f1_sc_pa, 'mcc_history': hist_mcc_sc_pa, 'rows': i_
16     {'Model': 'HardSampling Passive Aggressive Classifier', 'f1_history': hist_f1_hs_pa, 'mcc_history': hist_mcc_hs_pa, 'r_
17     {'Model': 'AdaBoost Logistic Regression', 'f1_history': hist_f1_ada, 'mcc_history': hist_mcc_ada, 'rows': rows_ada},_
18     {'Model': 'AdaBoost Hoeffding Tree', 'f1_history': hist_f1_ada_ht, 'mcc_history': hist_mcc_ada_ht, 'rows': rows_ada_ht}_
19     {'Model': 'AdaBoost PassiveAgressive Classifier', 'f1_history': hist_f1_ada_pa, 'mcc_history': hist_mcc_ada_pa, 'rows'_
20     {'Model': 'ADWINBoosting Logistic Regression', 'f1_history': hist_f1_adwin_lr, 'mcc_history': hist_mcc_adwin_lr, 'row:_
21     {'Model': 'ADWINBoosting PassiveAggressive Classifier', 'f1_history': hist_f1_adwin_pa, 'mcc_history': hist_mcc_adwin_p_
22 }]
23
24 # Plotting Historical F1 Scores
25 plt.figure(figsize=(15, 10))
26 for data in historical_data:
27     plt.plot(data['rows'], data['f1_history'], label=data['Model'], marker='o', markersize=5)
28 plt.title('Historical F1 Scores Over Time')
29 plt.xlabel('Number of Rows Processed')
30 plt.ylabel('F1 Score')
31 plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
32 plt.grid(True)
33 plt.tight_layout()
34 plt.show()
35

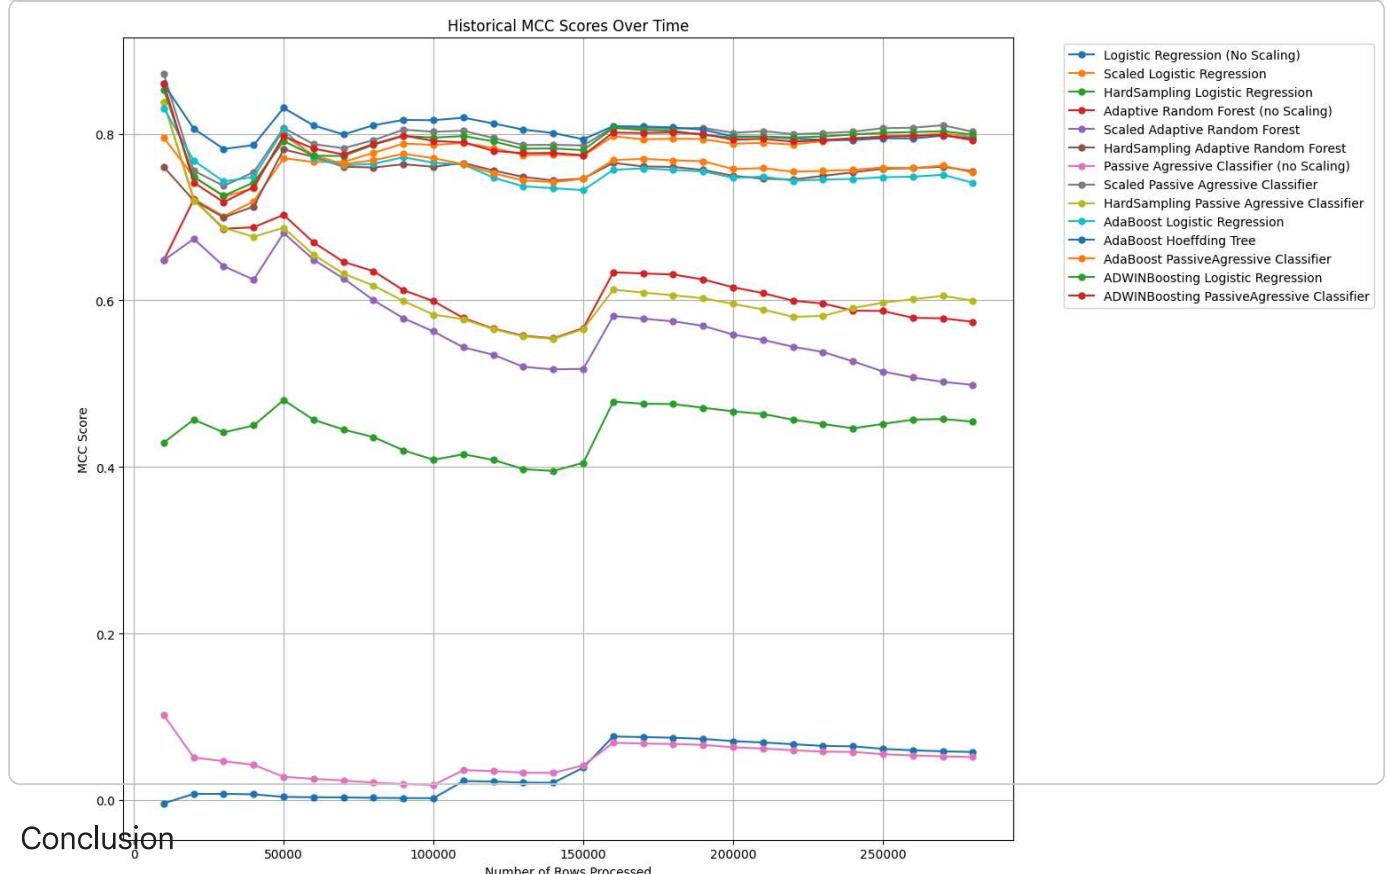
```



```

1 # Plotting Historical MCC Scores
2 plt.figure(figsize=(15, 10))
3 for data in historical_data:
4     plt.plot(data['rows'], data['mcc_history'], label=data['Model'],
5               marker='o', markersize=5)
5 plt.title('Historical MCC Scores Over Time')
6 plt.xlabel('Number of Rows Processed')
7 plt.ylabel('MCC Score')
8 plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
9 plt.grid(True)
10 plt.tight_layout()
11 plt.show()

```



## Conclusion

This notebook systematically evaluated various online machine learning models from the `River` library for credit card fraud detection. The models were compared based on their final F1 score, MCC score, and total runtime, with historical performance also visualized.

### Key Findings:

- **Impact of Scaling:** Feature scaling (e.g., using `StandardScaler`) had a significant positive impact on the performance of linear models like `LogisticRegression` and `PAClassifier`. Models without scaling often performed poorly, highlighting the importance of preprocessing for these algorithms in an online learning context.
- **Impact of HardSampling:** While `HardSamplingClassifier` was intended to address class imbalance, its effect on overall performance (F1 and MCC) varied. For Logistic Regression, HardSampling drastically reduced performance compared to its scaled counterpart. For `ARFCclassifier` and `PAClassifier`, HardSampling generally improved performance, indicating its utility can be model-specific and might require careful tuning.
- **Ensemble Methods:** Ensemble methods like `AdaBoostClassifier` and `ADWINBoostingClassifier` generally showed strong performance. `ADWINBoostingClassifier` with `LogisticRegression` as a base estimator achieved one of the highest F1 and MCC scores, demonstrating its robustness and adaptability to data streams.
- **Performance vs. Runtime:** There was often a trade-off between model complexity, performance, and runtime. Simple scaled models like `Scaled Passive Aggressive Classifier` and `Scaled Logistic Regression` offered competitive performance with relatively lower runtimes compared to more complex ensemble models like `AdaBoost Hoeffding Tree`.
- **Top Performing Models:**
  - `Scaled Logistic Regression` and `Scaled Passive Aggressive Classifier` consistently delivered high F1 and MCC

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.